A practical guide to building agents

Contents

Practical guide to building agents
Introduction
Large language models are becoming increasingly capable of handling complex, multi-step tas
Advances in reasoning, multimodality, and tool use have unlocked a new category of LLM-pow
systems known as agents.

This guide is designed for product and engineering teams exploring how to build their first ager
distilling insights from numerous customer deployments into practical and actionable best
practices. It includes frameworks for identifying promising use cases, clear patterns for designi
agent logic and orchestration, and best practices to ensure your agents run safely, predictably,
and effectively.

After reading this guide, you'll have the foundational knowledge you need to confidently start
building your first agent.

What is an agent?
While conventional software enables users to streamline and automate workflows, agents are a
to perform the same workflows on the users' behalf with a high degree of independence.

Agents are systems that independently accomplish tasks on your behalf.

A workflow is a sequence of steps that must be executed to meet the user's goal, whether that'
resolving a customer service issue, booking a restaurant reservation, committing a code chang
or generating a report.

Applications that integrate LLMs but don't use them to control workflow execution—think simpl
chatbots, single-turn LLMs, or sentiment classifiers—are not agents.

More concretely, an agent possesses core characteristics that allow it to act reliably and
consistently on behalf of a user.